

Moving windows

There are a number of ways to apply a function in a moving window. Here I review a couple of ideas. I found that with low numbers of data points simple for loops are more than sufficient, but the pandas implementation is far easier and faster so should be used. If you have a lot of data, then it may be worth taking the time to broadcast to a numpy array.

Applying a function

The standard functions which are applied in a moving window are averages and variances/std. As a result pandas has a built in method to handle this. To be fair to all methods, we will test with a user-defined function: the [mean absolute deviation](http://en.wikipedia.org/wiki/Absolute_deviation#Mean_absolute_deviation_.28MAD.29_.28about_mean.29) (http://en.wikipedia.org/wiki/Absolute_deviation#Mean_absolute_deviation_.28MAD.29_.28about_mean.29).

Lets get started by coding this function and some test data

```
In [52]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd


%matplotlib inline

# The user defined function we want to apply in a moving window
mad = lambda x: np.fabs(x - x.mean()).mean()

# Some random data
N = 1000
t = np.linspace(0, 10, N)
x = np.random.randn(N).cumsum()

# The moving window parameters
window_length = 300
window_shift = 1

# Plot the data
plt.plot(t, x)
plt.show()
```




The for loop

The simplest, but slowest, method is obviously the good old for loop. Let's test this first

```
In [53]: def ForLoop(f):
moving_val = []
moving_time = []
for i in xrange(0, N-window_length, window_shift):
moving_val.append(f(x[i:i+window_length]))
moving_time.append(np.average(t[i:i+window_length]))
return moving_time, moving_val

plt.plot(*ForLoop(mad))
plt.show()

%timeit ForLoop(mad)
```



10 loops, best of 3: 40.2 ms per loop


Pandas rolling_apply

Next up is the inspiration for all this, `pandas.rolling_apply`. This is by far the easiest method since it can be implemented in one line. However, as far as I can see there is no way to set the `window_shift`

```
In [54]: out = pd.rolling_apply(x, window_length, mad, center=True)

plt.plot(t, out)
plt.show()

%timeit out = pd.rolling_apply(x, window_length, mad)
```



10 loops, best of 3: 26.2 ms per loop

Numpy broadcast to array

Next we will broadcast the 1D array to a 2D array, compute the function along the new axis. This will require some effort to rewrite the function so it handles the shapes correctly. For help in understanding how this is done I really recommend taking a look at [this scipy page \(http://docs.scipy.org/doc/numpy/user/basics.broadcasting.html\)](http://docs.scipy.org/doc/numpy/user/basics.broadcasting.html)

```
In [55]: def NumpyArray():


    mad_array = lambda x: np.fabs(x.T - x.mean(axis=1)).mean(axis=0)

    vert_idx_list = np.arange(0, N - window_length, window_shift)
    hori_idx_list = np.arange(window_length)
    A, B = np.meshgrid(hori_idx_list, vert_idx_list)
    idx_array = A + B
    x_array = x[idx_array]

    return t[vert_idx_list+int(window_length/2.), mad_array(x_array)]

plt.plot(*NumpyArray())
plt.show()

%timeit NumpyArray()
```



100 loops, best of 3: 11.4 ms per loop

Additional

There is also a rolling apply function proposed by [Erik Rigtorp](http://www.rigtorp.se/2011/01/01/rolling-statistics-numpy.html) (<http://www.rigtorp.se/2011/01/01/rolling-statistics-numpy.html>). I still don't really understand how this works, but there is useful discussions to be found [here](http://numpy-discussion.10968.n7.nabble.com/Rolling-window-moving-average-moving-std-and-more-td4744.html) (<http://numpy-discussion.10968.n7.nabble.com/Rolling-window-moving-average-moving-std-and-more-td4744.html>). I've not included because I could not get it to work. If you can see how to do this please let me know!

In []: